

customer-churn-prediction-1

August 24, 2023

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import plotly.graph_objects as go
from plotly.subplots import make_subplots
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn import metrics
from sklearn.metrics import roc_curve
from sklearn.metrics import recall_score, confusion_matrix, precision_score, \
    f1_score, accuracy_score, classification_report
```

```
[3]: #loading data
df=pd.read_excel("C:\\Users\\pavan\\Downloads\\customer.xlsx")
```

```
[4]: df
```

```
[4]:      CustomerID      Name  Age  Gender  Location \
0              1  Customer_1   63    Male  Los Angeles
```

1	2	Customer_2	62	Female	New York
2	3	Customer_3	24	Female	Los Angeles
3	4	Customer_4	36	Female	Miami
4	5	Customer_5	46	Female	Miami
...
99995	99996	Customer_99996	33	Male	Houston
99996	99997	Customer_99997	62	Female	New York
99997	99998	Customer_99998	64	Male	Chicago
99998	99999	Customer_99999	51	Female	New York
99999	100000	Customer_100000	27	Female	Los Angeles

	Subscription_Length_Months	Monthly_Bill	Total_Usage_GB	Churn
0	17	73.36	236	0
1	1	48.76	172	0
2	5	85.47	460	0
3	3	97.94	297	1
4	19	58.14	266	0
...
99995	23	55.13	226	1
99996	19	61.65	351	0
99997	17	96.11	251	1
99998	20	49.25	434	1
99999	19	76.57	173	1

[100000 rows x 9 columns]

```
[5]: df.head()
```

```
[5]: CustomerID      Name  Age  Gender  Location \
0         1  Customer_1  63    Male  Los Angeles
1         2  Customer_2  62   Female   New York
2         3  Customer_3  24   Female  Los Angeles
3         4  Customer_4  36   Female    Miami
4         5  Customer_5  46   Female    Miami
```

	Subscription_Length_Months	Monthly_Bill	Total_Usage_GB	Churn
0	17	73.36	236	0
1	1	48.76	172	0
2	5	85.47	460	0
3	3	97.94	297	1
4	19	58.14	266	0

```
[6]: df.shape
```

```
[6]: (100000, 9)
```

```
[7]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CustomerID                            100000 non-null  int64
1   Name                                  100000 non-null  object
2   Age                                   100000 non-null  int64
3   Gender                                100000 non-null  object
4   Location                              100000 non-null  object
5   Subscription_Length_Months            100000 non-null  int64
6   Monthly_Bill                          100000 non-null  float64
7   Total_Usage_GB                        100000 non-null  int64
8   Churn                                 100000 non-null  int64
dtypes: float64(1), int64(5), object(3)
memory usage: 6.9+ MB

```

```
[8]: df.columns.values
```

```
[8]: array(['CustomerID', 'Name', 'Age', 'Gender', 'Location',
          'Subscription_Length_Months', 'Monthly_Bill', 'Total_Usage_GB',
          'Churn'], dtype=object)
```

```
[9]: df.dtypes
```

```
[9]: CustomerID                int64
     Name                    object
     Age                     int64
     Gender                  object
     Location                object
     Subscription_Length_Months  int64
     Monthly_Bill            float64
     Total_Usage_GB          int64
     Churn                   int64
     dtype: object

```

```
[10]: df = df.drop(['CustomerID'], axis = 1)
      df.head()
```

```
[10]:
```

	Name	Age	Gender	Location	Subscription_Length_Months	\
0	Customer_1	63	Male	Los Angeles		17
1	Customer_2	62	Female	New York		1
2	Customer_3	24	Female	Los Angeles		5
3	Customer_4	36	Female	Miami		3
4	Customer_5	46	Female	Miami		19

	Monthly_Bill	Total_Usage_GB	Churn
--	--------------	----------------	-------

0	73.36	236	0
1	48.76	172	0
2	85.47	460	0
3	97.94	297	1
4	58.14	266	0

```
[11]: df['Total_Usage_GB'] = pd.to_numeric(df.Total_Usage_GB, errors='coerce')
df.isnull().sum()
```

```
[11]: Name          0
      Age          0
      Gender       0
      Location     0
      Subscription_Length_Months  0
      Monthly_Bill  0
      Total_Usage_GB  0
      Churn        0
      dtype: int64
```

```
[12]: df[np.isnan(df['Total_Usage_GB'])]
```

```
[12]: Empty DataFrame
      Columns: [Name, Age, Gender, Location, Subscription_Length_Months, Monthly_Bill,
      Total_Usage_GB, Churn]
      Index: []
```

```
[13]: df[df['Age'] == 0].index
```

```
[13]: Int64Index([], dtype='int64')
```

```
[14]: df.drop(labels=df[df['Age'] == 0].index, axis=0, inplace=True)
df[df['Age'] == 0].index
```

```
[14]: Int64Index([], dtype='int64')
```

```
[15]: df.fillna(df["Total_Usage_GB"].mean())
```

```
[15]:
```

	Name	Age	Gender	Location	Subscription_Length_Months	\
0	Customer_1	63	Male	Los Angeles		17
1	Customer_2	62	Female	New York		1
2	Customer_3	24	Female	Los Angeles		5
3	Customer_4	36	Female	Miami		3
4	Customer_5	46	Female	Miami		19
...	
99995	Customer_99996	33	Male	Houston		23
99996	Customer_99997	62	Female	New York		19
99997	Customer_99998	64	Male	Chicago		17

99998	Customer_99999	51	Female	New York	20
99999	Customer_100000	27	Female	Los Angeles	19

	Monthly_Bill	Total_Usage_GB	Churn
0	73.36	236	0
1	48.76	172	0
2	85.47	460	0
3	97.94	297	1
4	58.14	266	0
...
99995	55.13	226	1
99996	61.65	351	0
99997	96.11	251	1
99998	49.25	434	1
99999	76.57	173	1

[100000 rows x 8 columns]

```
[16]: df.isnull().sum()
```

```
[16]: Name          0
      Age          0
      Gender       0
      Location     0
      Subscription_Length_Months  0
      Monthly_Bill  0
      Total_Usage_GB  0
      Churn        0
      dtype: int64
```

```
[17]: df["Monthly_Bill"].describe(include=['object', 'bool'])
```

```
[17]: count    100000.000000
      mean      65.053197
      std       20.230696
      min       30.000000
      25%       47.540000
      50%       65.010000
      75%       82.640000
      max       100.000000
      Name: Monthly_Bill, dtype: float64
```

```
[18]: numerical_cols = ['Age', 'Monthly_Bill', 'Total_Usage_GB']
      df[numerical_cols].describe()
```

```
[18]:           Age  Monthly_Bill  Total_Usage_GB
count  100000.000000  100000.000000  100000.000000
```

mean	44.027020	65.053197	274.393650
std	15.280283	20.230696	130.463063
min	18.000000	30.000000	50.000000
25%	31.000000	47.540000	161.000000
50%	44.000000	65.010000	274.000000
75%	57.000000	82.640000	387.000000
max	70.000000	100.000000	500.000000

```
[19]: g_labels = ['Male', 'Female']
      c_labels = ['No', 'Yes']
      # Create subplots: use 'domain' type for Pie subplot
      fig = make_subplots(rows=1, cols=2, specs=[[{'type': 'domain'}, {'type':
      ↪ 'domain'}]])
      fig.add_trace(go.Pie(labels=g_labels, values=df['Gender'].value_counts(),
      ↪ name="Gender"),
                    1, 1)
      fig.add_trace(go.Pie(labels=c_labels, values=df['Age'].value_counts(),
      ↪ name="Age"),
                    1, 2)

      # Use `hole` to create a donut-like pie chart
      fig.update_traces(hole=.4, hoverinfo="label+percent+name", textfont_size=16)

      fig.update_layout(
          title_text="customer_churn_prediction",
          # Add annotations in the center of the donut pies.
          annotations=[dict(text='Gender', x=0.16, y=0.5, font_size=20,
      ↪ showarrow=False),
                       dict(text='Age', x=0.84, y=0.5, font_size=20,
      ↪ showarrow=False)])
      fig.show()
```

```
[20]: df["Gender"][df["Gender"]=="No"].groupby(by=df["Age"]).count()
```

```
[20]: Series([], Name: Gender, dtype: int64)
```

```
[21]: df["Age"][df["Age"]=="Yes"].groupby(by=df["Gender"]).count()
```

```
[21]: Series([], Name: Age, dtype: int64)
```

```
[22]: plt.figure(figsize=(6, 6))
      labels = ["Churn: Yes", "Churn: No"]
      values = [1869, 5163]
      labels_gender = ["F", "M", "F", "M"]
      sizes_gender = [939, 930, 2544, 2619]
      colors = ['#ff6666', '#66b3ff']
      colors_gender = ['#c2c2f0', '#ffb3e6', '#c2c2f0', '#ffb3e6']
```

```

explode = (0.3,0.3)
explode_gender = (0.1,0.1,0.1,0.1)
textprops = {"fontsize":15}
#Plot
plt.pie(values, labels=labels,autopct='%1.1f%%',pctdistance=1.08,
    ↪labeldistance=0.8,colors=colors, startangle=90,frame=True,
    ↪explode=explode,radius=10, textprops =textprops, counterclock = True, )
plt.pie(sizes_gender,labels=labels_gender,colors=colors_gender,startangle=90,
    ↪explode=explode_gender,radius=7, textprops =textprops, counterclock = True, )
#Draw circle
centre_circle = plt.Circle((0,0),5,color='black', fc='white',linewidth=0)
fig = plt.gcf()
fig.gca().add_artist(centre_circle)


plt.title('Churn Distribution w.r.t Gender: Male(M), Female(F)', fontsize=15,
    ↪y=1.1)

# show plot

plt.axis('equal')
plt.tight_layout()
plt.show()

```

Churn Distribution w.r.t Gender: Male(M), Female(F)



```
[24]: labels = df['Monthly_Bill'].unique()
      values = df['Monthly_Bill'].value_counts()

      fig = go.Figure(data=[go.Pie(labels=labels, values=values, hole=.3)])
      fig.update_layout(title_text="<b>Payment Method Distribution</b>")
      fig.show()
```

```
[25]: fig = px.histogram(df, x="Churn", color="Monthly_Bill", title="<b>Customer_
      ↳Payment Method distribution w.r.t. Churn</b>")
      fig.update_layout(width=700, height=500, bargap=0.1)
      fig.show()
```



```
[26]: df["Subscription_Length_Months"].unique()
```

```
[26]: array([17,  1,  5,  3, 19, 15, 10, 12, 20, 13,  8, 23,  2,  4, 18,  9, 14,  
        16,  6,  7, 24, 22, 11, 21], dtype=int64)
```

```
[27]: df[df["Gender"]=="Male"][["Subscription_Length_Months", "Churn"]].value_counts()
```

```
[27]: Subscription_Length_Months  Churn  
1                                1      1109  
2                                1      1089  
14                               0      1088  
21                               0      1088  
20                               0      1087  
22                               0      1083  
15                               0      1079  
22                               1      1067  
10                               1      1067  
2                                0      1059  
20                               1      1058  
13                               1      1058  
5                                0      1056  
14                               1      1056  
11                               1      1056  
                                0      1050  
16                               1      1049  
18                               1      1046  
7                                1      1044  
9                                0      1042  
10                               0      1040  
12                               1      1040  
16                               0      1038  
3                                1      1037  
24                               0      1037  
17                               0      1034  
7                                0      1034  
4                                0      1030  
8                                1      1027  
19                               1      1025  
9                                1      1024  
5                                1      1024  
24                               1      1024  
8                                0      1022  
6                                0      1020  
18                               0      1018  
23                               0      1016  
19                               0      1012  
1                                0      1011
```

3	0	1010
23	1	1006
12	0	1001
6	1	1000
4	1	1000
13	0	994
15	1	980
21	1	977
17	1	972

dtype: int64

```
[28]: df[df["Gender"]=="Female"][["Subscription_Length_Months", "Churn"]].
      ↪value_counts()
```

```
[28]: Subscription_Length_Months  Churn
6                                0      1124
5                                0      1099
1                                0      1098
20                               0      1081
18                               0      1081
16                               1      1080
20                               1      1077
22                               1      1072
7                                1      1071
12                               0      1069
16                               0      1062
7                                0      1062
21                               1      1061
2                                0      1060
24                               1      1059
11                               0      1058
9                                0      1055
13                               0      1053
14                               0      1052
4                                0      1050
3                                1      1049
17                               1      1049
13                               1      1049
19                               1      1047
12                               1      1045
22                               0      1045
8                                0      1044
23                               1      1042
6                                1      1040
3                                0      1040
11                               1      1036
10                               0      1034
```

15	0	1034
	1	1029
1	1	1029
21	0	1028
18	1	1026
19	0	1022
2	1	1020
23	0	1019
4	1	1018
14	1	1017
8	1	1013
9	1	1013
10	1	1010
17	0	1009
24	0	993
5	1	992

dtype: int64

```
[29]: fig = go.Figure()

fig.add_trace(go.Bar(
    x = [['Churn:No', 'Churn:No', 'Churn:Yes', 'Churn:Yes'],
         ["Female", "Male", "Female", "Male"]],
    y = [965, 992, 219, 240],
    name = 'DSL',
))

fig.add_trace(go.Bar(
    x = [['Churn:No', 'Churn:No', 'Churn:Yes', 'Churn:Yes'],
         ["Female", "Male", "Female", "Male"]],
    y = [889, 910, 664, 633],
    name = 'Fiber optic',
))

fig.add_trace(go.Bar(
    x = [['Churn:No', 'Churn:No', 'Churn:Yes', 'Churn:Yes'],
         ["Female", "Male", "Female", "Male"]],
    y = [690, 717, 56, 57],
    name = 'No Internet',
))

fig.update_layout(title_text="<b>Churn Distribution w.r.t. Internet Service and_
↳Gender</b>")

fig.show()
```

```
[30]: color_map = {"Yes": "#FF97FF", "No": "#AB63FA"}
fig = px.histogram(df, x="Churn", color="Subscription_Length_Months",
    ↪barmode="group", title="<b>Dependents distribution</b>",
    ↪color_discrete_map=color_map)
fig.update_layout(width=700, height=500, bargap=0.1)
fig.show()
```

```
[31]: color_map = {"Yes": '#FFA15A', "No": '#00CC96'}
fig = px.histogram(df, x="Churn", color="Monthly_Bill", barmode="group",
    ↪title="<b>Chrun distribution w.r.t. Partners</b>",
    ↪color_discrete_map=color_map)
fig.update_layout(width=700, height=500, bargap=0.1)
fig.show()
```

```
[32]: color_map = {"Yes": '#00CC96', "No": '#B6E880'}
fig = px.histogram(df, x="Churn", color="Age", title="<b>Chrun distribution w.r.
    ↪t. Senior Citizen</b>", color_discrete_map=color_map)
fig.update_layout(width=700, height=500, bargap=0.1)
fig.show()
```

```
[33]: color_map = {"Yes": "#FF97FF", "No": "#AB63FA"}
fig = px.histogram(df, x="Churn", color="Age", barmode="group", title="<b>Churn
    ↪w.r.t Online Security</b>", color_discrete_map=color_map)
fig.update_layout(width=700, height=500, bargap=0.1)
fig.show()
```

```
[34]: color_map = {"Yes": '#FFA15A', "No": '#00CC96'}
fig = px.histogram(df, x="Churn", color="Total_Usage_GB", title="<b>Chrun
    ↪distribution w.r.t. Paperless Billing</b>", color_discrete_map=color_map)
fig.update_layout(width=700, height=500, bargap=0.1)
fig.show()
```

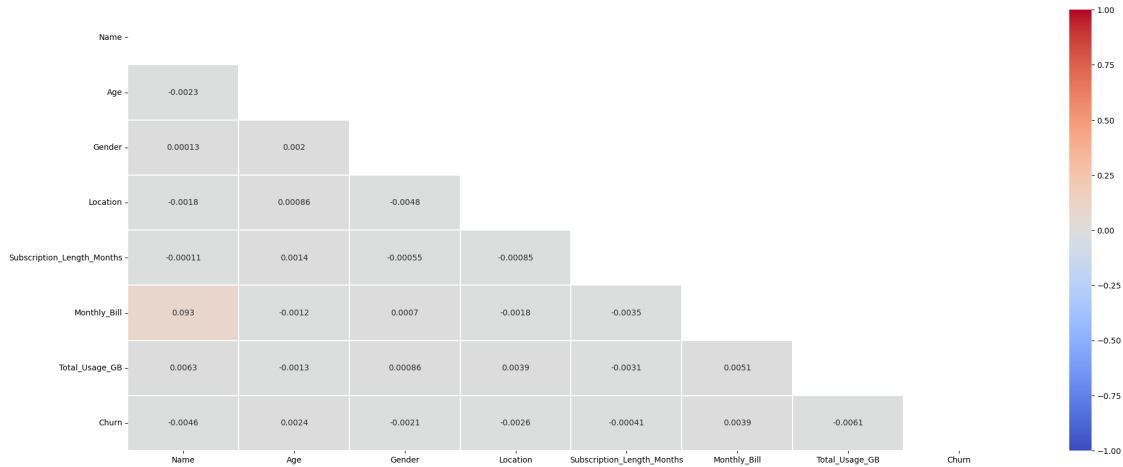
```
[35]: fig = px.histogram(df, x="Churn", color="Gender",barmode="group",
    ↪title="<b>Chrun distribution w.r.t. TechSupport</b>")
fig.update_layout(width=700, height=500, bargap=0.1)
fig.show()
```

```
[36]: plt.figure(figsize=(25, 10))

corr = df.apply(lambda x: pd.factorize(x)[0]).corr()

mask = np.triu(np.ones_like(corr, dtype=bool))

ax = sns.heatmap(corr, mask=mask, xticklabels=corr.columns, yticklabels=corr.
    ↪columns, annot=True, linewidths=.2, cmap='coolwarm', vmin=-1, vmax=1)
```



```
[37]: def object_to_int(dataframe_series):
        if dataframe_series.dtype=='object':
            dataframe_series = LabelEncoder().fit_transform(dataframe_series)
        return dataframe_series
```

```
[38]: df = df.apply(lambda x: object_to_int(x))
df.head()
```

```
[38]:
```

	Name	Age	Gender	Location	Subscription_Length_Months	Monthly_Bill	Total_Usage_GB	Churn
0	0	63	1	2	17	73.36	236	0
1	11112	62	0	4	1	48.76	172	0
2	22223	24	0	2	5	85.47	460	0
3	33334	36	0	3	3	97.94	297	1
4	44445	46	0	3	19	58.14	266	0

```
[39]: plt.figure(figsize=(14,7))
df.corr()['Churn'].sort_values(ascending = False)
```

```
[39]: Churn          1.000000
Location          0.006405
Subscription_Length_Months  0.002328
Gender            0.002121
Age              0.001559
Monthly_Bill     -0.000211
```

```
Name -0.001418
Total_Usage_GB -0.002842
Name: Churn, dtype: float64
```

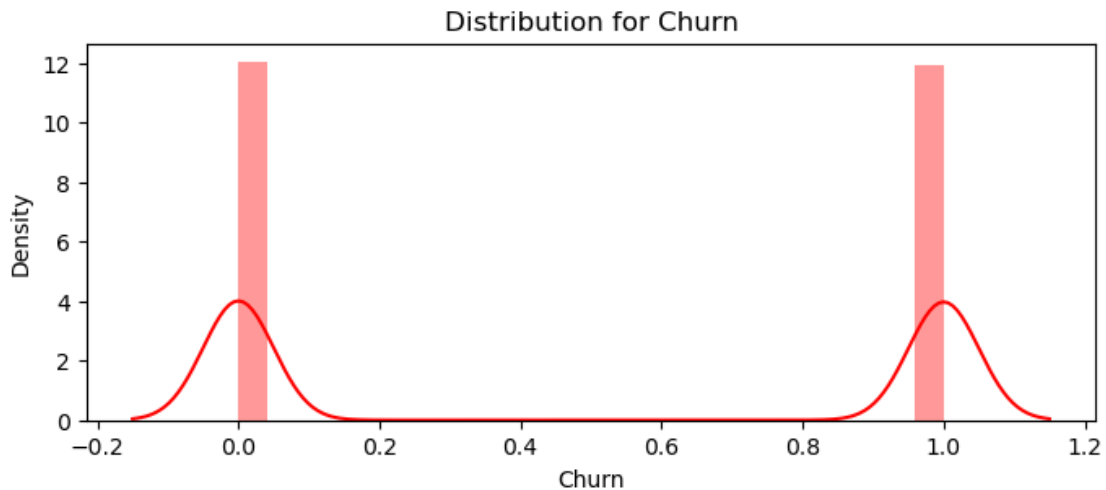
<Figure size 1400x700 with 0 Axes>



```
[40]: X = df.drop(columns = ['Churn'])
      y = df['Churn'].values
```

```
[41]: X_train, X_test, y_train, y_test = train_test_split(X,y,test_size = 0.30,
      ↪random_state = 40, stratify=y)
```

```
[42]: def distplot(feature, frame, color='r'):
      plt.figure(figsize=(8,3))
      plt.title("Distribution for {}".format(feature))
      ax = sns.distplot(frame[feature], color= color)
```

```
[43]: num_cols = ["Churn", 'Monthly_Bill', 'Total_Usage_GB']
      for feat in num_cols: distplot(feat, df)
```





```
[44]: knn_model = KNeighborsClassifier(n_neighbors = 11)
knn_model.fit(X_train,y_train)
predicted_y = knn_model.predict(X_test)
accuracy_knn = knn_model.score(X_test,y_test)
print("KNN accuracy:",accuracy_knn)
```

KNN accuracy: 0.49773333333333336

```
[45]: print(classification_report(y_test, predicted_y))
```

	precision	recall	f1-score	support
0	0.50	0.50	0.50	15066
1	0.50	0.49	0.49	14934

accuracy			0.50	30000
macro avg	0.50	0.50	0.50	30000
weighted avg	0.50	0.50	0.50	30000

```
[46]: svc_model = SVC(random_state = 1)
svc_model.fit(X_train,y_train)
predict_y = svc_model.predict(X_test)
accuracy_svc = svc_model.score(X_test,y_test)
print("SVM accuracy is :",accuracy_svc)
```

SVM accuracy is : 0.5022

```
[47]: print(classification_report(y_test, predict_y))
```

		precision	recall	f1-score	support
	0	0.50	1.00	0.67	15066
	1	0.00	0.00	0.00	14934
accuracy				0.50	30000
macro avg		0.25	0.50	0.33	30000
weighted avg		0.25	0.50	0.34	30000

```
[48]: model_rf = RandomForestClassifier(n_estimators=500 , oob_score = True, n_jobs = -1,
                                     random_state =50, max_features = "auto",
                                     max_leaf_nodes = 30)
model_rf.fit(X_train, y_train)

# Make predictions
prediction_test = model_rf.predict(X_test)
print (metrics.accuracy_score(y_test, prediction_test))
```


0.4997333333333333

```
[49]: print(classification_report(y_test, prediction_test))
```


		precision	recall	f1-score	support
	0	0.50	0.64	0.56	15066
	1	0.50	0.36	0.41	14934
accuracy				0.50	30000
macro avg		0.50	0.50	0.49	30000
weighted avg		0.50	0.50	0.49	30000


```
[50]: plt.figure(figsize=(4,3))
sns.heatmap(confusion_matrix(y_test, prediction_test),
            annot=True,fmt = "d",linecolor="k",linewidths=3)

plt.title(" RANDOM FOREST CONFUSION MATRIX",fontsize=14)
plt.show()
```



```
[51]: y_rfpred_prob = model_rf.predict_proba(X_test)[:,-1]
fpr_rf, tpr_rf, thresholds = roc_curve(y_test, y_rfpred_prob)
plt.plot([0, 1], [0, 1], 'k--' )
plt.plot(fpr_rf, tpr_rf, label='Random Forest',color = "r")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Random Forest ROC Curve',fontsize=16)
plt.show();
```



```
[52]: lr_model = LogisticRegression()
lr_model.fit(X_train,y_train)
accuracy_lr = lr_model.score(X_test,y_test)
print("Logistic Regression accuracy is :",accuracy_lr)
```

Logistic Regression accuracy is : 0.5022

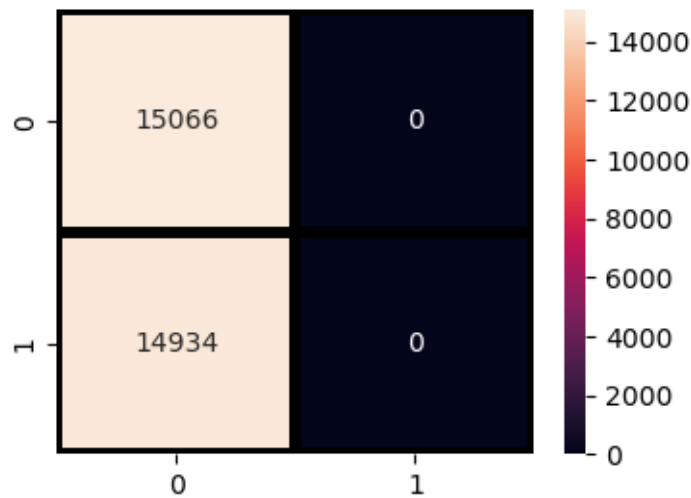
```
[53]: lr_pred= lr_model.predict(X_test)
report = classification_report(y_test,lr_pred)
print(report)
```

	precision	recall	f1-score	support
0	0.50	1.00	0.67	15066
1	0.00	0.00	0.00	14934
accuracy			0.50	30000
macro avg	0.25	0.50	0.33	30000
weighted avg	0.25	0.50	0.34	30000


```
[54]: plt.figure(figsize=(4,3))
sns.heatmap(confusion_matrix(y_test, lr_pred),
            annot=True,fmt = "d",linecolor="k",linewidths=3)

plt.title("LOGISTIC REGRESSION CONFUSION MATRIX",fontsize=14)
plt.show()
```

LOGISTIC REGRESSION CONFUSION MATRIX



```
[55]: y_pred_prob = lr_model.predict_proba(X_test)[:,-1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
plt.plot([0, 1], [0, 1], 'k--' )
plt.plot(fpr, tpr, label='Logistic Regression',color = "r")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Logistic Regression ROC Curve',fontsize=16)
plt.show();
```



```
[56]: dt_model = DecisionTreeClassifier()
dt_model.fit(X_train,y_train)
predictdt_y = dt_model.predict(X_test)
accuracy_dt = dt_model.score(X_test,y_test)
print("Decision Tree accuracy is :",accuracy_dt)
```

Decision Tree accuracy is : 0.4982

```
[57]: print(classification_report(y_test, predictdt_y))
```

	precision	recall	f1-score	support
0	0.50	0.50	0.50	15066
1	0.50	0.49	0.50	14934
accuracy			0.50	30000
macro avg	0.50	0.50	0.50	30000
weighted avg	0.50	0.50	0.50	30000

```
[58]: a_model = AdaBoostClassifier()
a_model.fit(X_train,y_train)
a_preds = a_model.predict(X_test)
print("AdaBoost Classifier accuracy")
metrics.accuracy_score(y_test, a_preds)
```

AdaBoost Classifier accuracy


```
[58]: 0.4981333333333333
```

```
[59]: print(classification_report(y_test, a_preds))
```

	precision	recall	f1-score	support
0	0.50	0.58	0.54	15066
1	0.50	0.41	0.45	14934
accuracy			0.50	30000
macro avg	0.50	0.50	0.49	30000
weighted avg	0.50	0.50	0.49	30000

```
[60]: plt.figure(figsize=(4,3))
sns.heatmap(confusion_matrix(y_test, a_preds),
            annot=True,fmt = "d",linecolor="k",linewidths=3)

plt.title("AdaBoost Classifier Confusion Matrix",fontsize=14)
plt.show()
```



```
[61]: gb = GradientBoostingClassifier()
gb.fit(X_train, y_train)
gb_pred = gb.predict(X_test)
print("Gradient Boosting Classifier", accuracy_score(y_test, gb_pred))
```

Gradient Boosting Classifier 0.4958666666666667


```
[62]: print(classification_report(y_test, gb_pred))
```

	precision	recall	f1-score	support
0	0.50	0.55	0.52	15066
1	0.49	0.44	0.47	14934
accuracy			0.50	30000
macro avg	0.50	0.50	0.49	30000
weighted avg	0.50	0.50	0.49	30000

```
[63]: plt.figure(figsize=(4,3))
sns.heatmap(confusion_matrix(y_test, gb_pred),
            annot=True,fmt = "d",linecolor="k",linewidths=3)

plt.title("Gradient Boosting Classifier Confusion Matrix",fontsize=14)
plt.show()
```

Gradient Boosting Classifier Confusion Matrix



```
[64]: from sklearn.ensemble import VotingClassifier
clf1 = GradientBoostingClassifier()
```

```

clf2 = LogisticRegression()
clf3 = AdaBoostClassifier()
ecf1 = VotingClassifier(estimators=[('gbc', clf1), ('lr', clf2), ('abc',
    ↪clf3)], voting='soft')
ecf1.fit(X_train, y_train)
predictions = ecf1.predict(X_test)
print("Final Accuracy Score ")
print(accuracy_score(y_test, predictions))

```

Final Accuracy Score
0.4968

```
[65]: print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.50	0.61	0.55	15066
1	0.49	0.38	0.43	14934
accuracy			0.50	30000
macro avg	0.50	0.50	0.49	30000
weighted avg	0.50	0.50	0.49	30000

```

[66]: plt.figure(figsize=(4,3))
sns.heatmap(confusion_matrix(y_test, predictions),
            annot=True,fmt = "d",linecolor="k",linewidths=3)

plt.title("FINAL CONFUSION MATRIX",fontsize=14)
plt.show()

```

